

[12pt,a4paper]article amsfonts amsmath amssymb geometry indentfirst left=2.0cm,right=2.0cm,top=2.5cm

# Markov Decision Processes (MDP)

esgl Hu

March 12, 2018

## 1 Background

aaaaaaaaa

## 2 Notation & Definition

A reinforcement learning task that satisfies the Markov property is called a *Markov decision process*, or *MDP*. If the state and action spaces are finite, then it is called a *finite Markov decision process* (*finite MDP*)

Given any state and action  $s$  and  $a$ , the probability of each possible pair of next state and reward,  $s', r$ , is denoted

$$p(s', r|s, a) \doteq \Pr(\{S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a\}) \quad (1)$$

Given the dynamics as specified by (1), one can compute anything else one might want to know about the environment, such as the expected rewards for state-action pairs.

$$r(s, a) \doteq \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r|s, a), \quad (2)$$

the *state-transition probabilities*,

$$p(s'|s, a) \doteq \Pr\{S_{t+1} = s' | S_t = s, A_t = a\} = \sum_{r \in \mathcal{R}} p(s', r|s, a) \quad (3)$$

and the expected rewards for state-action-next-state triples,

$$r(s, a, s') \doteq \mathbb{E}[R_{t+1} | S_t = s, A_t = a, S_{t+1} = s'] = \frac{\sum_{r \in \mathcal{R}} r p(s', r|s, a)}{p(s'|s, a)} \quad (4)$$

### 3 preface

#### Notations & Definitions

The methods have learned the values of actions and then selected actions based on their estimated action values, their policies would not even exist without the action-value estimates.

The methods have learned a *parameterized policy* that can select actions without consulting a value function. A value function may still be used to *learn* the policy weights, but is not required for action selection.

The notation  $\theta \in \mathbb{R}^n$  is used for the primary learned weight vector,  $\pi(a|s, \theta) \doteq \Pr\{A_t = a|S_t = s, \theta_t = \theta\}$  for the probability that action  $a$  is taken at time  $t$  given that the agent is in state  $s$  at time  $t$  weight vector  $\theta$ .

These methods seek to maximize performance  $\eta(\theta)$ , so their updates approximate gradient ascent in  $\eta$ :

$$\theta_{t+1} \doteq \theta_t + \alpha \widehat{\nabla(\eta(\theta_t))} \quad (5)$$

where  $\widehat{\nabla(\eta(\theta_t))}$  is a stochastic estimate whose expectation approximates the gradients of the performance measure  $\eta(\theta_t)$  with respect to its argument  $\theta_t$

#### Advantages & Disadvantages

##### Advantages

- Better convergence properties
- Effective in high-dimensional or continuous action spaces
- Can learn stochastic policies

##### Disadvantages

- Typically converge to a local rather than global optimum
- Evaluating a policy is typically inefficient and high variance

#### Definition of Policy Gradient

We let  $\tau$  denote a state-action sequence  $\{s_0, a_0, \dots, s_H, a_H\}$ . We overload notation:  $R(\tau) = \sum_{t=0}^H R(s_t, a_t)$ .

$$\eta(\theta) = \mathbb{E}\left[\sum_{t=0}^H R(s_t, a_t), \pi_\theta\right] = \sum_{\tau} P(\tau; \theta) R(\tau) \quad (6)$$

In our new notation, our goal is to find  $\theta$ :

$$\max_{\theta} \eta(\theta) = \max_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau) \quad (7)$$

REINFORCE: A Monte-Carlo Policy-Gradient Method (episodic) a differentiable policy parameterization  $\pi(a|s, \theta)$ ,  $\forall a \in A, s \in S, \theta \in R^n$  Initialize policy weights  $\theta$  True Generate an episode  $\{S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T, S_T\}$  follow  $\pi(\cdot|\cdot, \theta)$  each step of the episode  $t = 0, 1, \dots, T-1$   $G_t \leftarrow$  return from step  $t$   $\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_{\theta} \log \pi(A_t|S_t, \theta)$